

# Intelligent Systems

## Assignment 2

Bernardo Morais Chagas (103639)

September 30, 2025

Repository link: [\*\*GitHub Repository\*\*](#)

```
In [1]: import numpy as np
        from sklearn import datasets
        from sklearn.preprocessing import StandardScaler
        from sklearn.model_selection import train_test_split
        from sklearn.metrics import mean_squared_error, accuracy_score, classification_report
        import skfuzzy as fuzz
        import matplotlib.pyplot as plt
        import torch
        import torch.nn as nn
        import torch.optim as optim
        import pandas
```

Importation of Dataset 1

```
In [2]: # CHOOSE DATASET

        # Binary classification dataset
        diabetes = datasets.load_diabetes(as_frame=True)

        # Regression dataset
        #data = datasets.fetch_openml(name="boston", version=1, as_frame=True)

        X = diabetes.data.values
        y = diabetes.target.values

        print("Shape:", X.shape)

        print(diabetes.data.head(), "\n \n") # first rows of features
        print(diabetes.target.head()) # first rows of target
```

Shape: (442, 10)

	age	sex	bmi	bp	s1	s2	s3	\
0	0.038076	0.050680	0.061696	0.021872	-0.044223	-0.034821	-0.043401	
1	-0.001882	-0.044642	-0.051474	-0.026328	-0.008449	-0.019163	0.074412	
2	0.085299	0.050680	0.044451	-0.005670	-0.045599	-0.034194	-0.032356	
3	-0.089063	-0.044642	-0.011595	-0.036656	0.012191	0.024991	-0.036038	
4	0.005383	-0.044642	-0.036385	0.021872	0.003935	0.015596	0.008142	

	s4	s5	s6
0	-0.002592	0.019907	-0.017646
1	-0.039493	-0.068332	-0.092204
2	-0.002592	0.002861	-0.025930
3	0.034309	0.022688	-0.009362
4	-0.002592	-0.031988	-0.046641

0	151.0
1	75.0
2	141.0
3	206.0
4	135.0

Name: target, dtype: float64

```
In [3]: #train test splitting
test_size=0.2
Xtr, Xte, ytr, yte = train_test_split(X, y, test_size=test_size, random_state=42)
```

```
In [4]: # Standardize features
scaler=StandardScaler()
Xtr= scaler.fit_transform(Xtr)
Xte= scaler.transform(Xte)
```

In order to be able to compare the results, the number of clusters and value of m used, was the same as the one in the previous assignig ( n\_clusters = 4; m=1.1)

```
In [5]: # Number of clusters
n_clusters = 4
m=1.1
```

```

# Concatenate target for clustering
Xexp=np.concatenate([Xtr, ytr.reshape(-1, 1)], axis=1)
#Xexp=Xtr

# Transpose data for skfuzzy (expects features x samples)
Xexp_T = Xexp.T

# Fuzzy C-means clustering
centers, u, u0, d, jm, p, fpc = fuzz.cluster.cmeans(
    Xexp_T, n_clusters, m=m, error=0.005, maxiter=1000, init=None,
)

```

In [6]: centers.shape

Out[6]: (4, 11)

```

In [7]: # Compute sigma (spread) for each cluster
sigmas = []
for j in range(n_clusters):
    # membership weights for cluster j, raised to m
    u_j = u[j, :] ** m
    # weighted variance for each feature
    var_j = np.average((Xexp - centers[j])**2, axis=0, weights=u_j)
    sigma_j = np.sqrt(var_j)
    sigmas.append(sigma_j)
sigmas=np.array(sigmas)

```

```


In [8]: # Hard clustering from fuzzy membership
cluster_labels = np.argmax(u, axis=0)
print("Fuzzy partition coefficient (FPC):", fpc)

# Plot first two features with fuzzy membership
plt.figure(figsize=(8,6))
for j in range(n_clusters):
    plt.scatter(
        Xexp[cluster_labels == j, 0],          # Feature 1
        Xexp[cluster_labels == j, 1],          # Feature 2
        alpha=u[j, :],                         # transparency ~ membership
        label=f'Cluster {j}'
    )

```


```
plt.title("Fuzzy C-Means Clustering (with membership degree)")  
plt.xlabel("Feature 1")  
plt.ylabel("Feature 2")  
plt.legend()  
plt.show()
```

Fuzzy partition coefficient (FPC): 0.982771992053368



```
In [9]: # Plot first two features with cluster assignments
plt.figure(figsize=(8,6))
for j in range(n_clusters):
    plt.scatter(
        Xexp[cluster_labels == j, 0],
        Xexp[cluster_labels == j, 1],
        label=f'Cluster {j}'
    )

plt.title("Fuzzy C-Means Clustering (CRISPEN)")
plt.xlabel("Feature 1")
plt.ylabel("Feature 2")
plt.legend()
plt.show()
```



```
In [10]: # Gaussian formula
def gaussian(x, mu, sigma):
    return np.exp(-0.5 * ((x - mu)/sigma)**2)


lin=np.linspace(-2, 4, 500)
plt.figure(figsize=(8,6))
```

```
y_aux=[]
feature=0
for j in range(n_clusters):
    # Compute curves
    y_aux.append(gaussian(lin, centers[j,feature], sigmas[j,feature]))

    # Plot
    plt.plot(lin, y_aux[j], label=f"Gaussian  $\mu=\{\text{np.round}(\text{centers}[j,\text{feature}],2)\}$ ,  $\sigma=\{\text{np.round}(\text{sigmas}[j,\text{feature}],2)\}$ ")

plt.title("Projection of the membership functions on Feature 2")
plt.xlabel("Feature 1")
plt.ylabel("Degree of Membership")
plt.legend()
plt.grid(True)
plt.show()
```





```
In [11]: # -----
# Gaussian Membership Function
# -----
class GaussianMF(nn.Module):
    def __init__(self, centers, sigmas, agg_prob):
        super().__init__()
        self.centers = nn.Parameter(torch.tensor(centers, dtype=torch.float32))
```

```

self.sigmas = nn.Parameter(torch.tensor(sigmas, dtype=torch.float32))
self.agg_prob=agg_prob

def forward(self, x):
    # Expand for broadcasting
    # x: (batch, 1, n_dims), centers: (1, n_rules, n_dims), sigmas: (1, n_rules, n_dims)
    diff = abs((x.unsqueeze(1) - self.centers.unsqueeze(0))/self.sigmas.unsqueeze(0)) #(batch, n_rules, n_dims)

    # Aggregation
    if self.agg_prob:
        dist = torch.norm(diff, dim=-1) # (batch, n_rules) # probabilistic intersection
    else:
        dist = torch.max(diff, dim=-1).values # (batch, n_rules) # min intersection (min instersection of normal funtion

    return torch.exp(-0.5 * dist ** 2)

# -----
# TSK Model
# -----
class TSK(nn.Module):
    def __init__(self, n_inputs, n_rules, centers, sigmas,agg_prob=False):
        super().__init__()
        self.n_inputs = n_inputs
        self.n_rules = n_rules

        # Antecedents (Gaussian MFs)

        self.mfs=GaussianMF(centers, sigmas,agg_prob)

        # Consequents (Linear functions of inputs)
        # Each rule has coeffs for each input + bias
        self.consequents = nn.Parameter(
            torch.randn(n_inputs + 1,n_rules)
        )

    def forward(self, x):
        # x: (batch, n_inputs)
        batch_size = x.shape[0]

        # Compute membership values for each input feature

```

```

# firing_strengths: (batch, n_rules)
firing_strengths = self.mfs(x)

# Normalize memberships
# norm_fs: (batch, n_rules)
norm_fs = firing_strengths / (firing_strengths.sum(dim=1, keepdim=True) + 1e-9)

# Consequent output (linear model per rule)
x_aug = torch.cat([x, torch.ones(batch_size, 1)], dim=1) # add bias

rule_outputs = torch.einsum("br,rk->bk", x_aug, self.consequents) # (batch, rules)
# Weighted sum
output = torch.sum(norm_fs * rule_outputs, dim=1, keepdim=True)

return output, norm_fs, rule_outputs

```

```

In [12]: # -----
# Least Squares Solver for Consequents (TSK)
# -----
def train_ls(model, X, y):
    with torch.no_grad():
        _, norm_fs, _ = model(X)

        # Design matrix for LS: combine normalized firing strengths with input
        X_aug = torch.cat([X, torch.ones(X.shape[0], 1)], dim=1)

        Phi = torch.einsum("br,bi->bri", X_aug, norm_fs).reshape(X.shape[0], -1)

        # Solve LS: consequents = (Phi^T Phi)^-1 Phi^T y

        theta = torch.linalg.lstsq(Phi, y).solution

        model.consequents.data = theta.reshape(model.consequents.shape)

```

```

In [13]: # -----
# Gradient Descent Training
# -----

```

```
def train_gd(model, X, y, epochs=100, lr=1e-3):
    optimizer = optim.Adam(model.parameters(), lr=lr)
    criterion = nn.MSELoss()
    for _ in range(epochs):
        optimizer.zero_grad()
        y_pred, _, _ = model(X)
        loss = criterion(y_pred, y)
        print(loss)
        loss.backward()
        optimizer.step()
```

```
In [14]: # -----
# Hybrid Training (Classic ANFIS)
# -----
def train_hybrid_anfis(model, X, y, max_iters=10, gd_epochs=20, lr=1e-3): #10, 20, 1e-3
    train_ls(model, X, y)
    for _ in range(max_iters):
        # Step A: GD on antecedents (freeze consequents)
        model.consequents.requires_grad = False
        train_gd(model, X, y, epochs=gd_epochs, lr=lr)

        # Step B: LS on consequents (freeze antecedents)
        model.consequents.requires_grad = True
        model.mfs.requires_grad = False
        train_ls(model, X, y)

        # Re-enable antecedents
        model.mfs.requires_grad = True
```

```
In [15]: # -----
# Alternative Hybrid Training (LS+ gradient descent on all)
# -----
def train_hybrid(model, X, y, epochs=100, lr=1e-5): #def; 100, 4
    # Step 1: LS for consequents
    train_ls(model, X, y)
    # Step 2: GD fine-tuning
    train_gd(model, X, y, epochs=epochs, lr=lr)
```

```
In [16]: # Build model
model = TSK(n_inputs=Xtr.shape[1], n_rules=n_clusters, centers=centers[:, :-1], sigmas=sigmas[:, :-1])
```

```
Xtr = torch.tensor(Xtr, dtype=torch.float32)
ytr = torch.tensor(ytr, dtype=torch.float32)
Xte = torch.tensor(Xte, dtype=torch.float32)
yte = torch.tensor(yte, dtype=torch.float32)
```

```
In [17]: # Training with LS:
#train_ls(model, Xtr, ytr.reshape(-1,1))
train_hybrid_anfis(model, Xtr, ytr.reshape(-1,1), max_iters=10, gd_epochs=14, lr=1e-4) #10 20 3
```

tensor(2382.7402, grad\_fn=<MseLossBackward0>)  
tensor(2382.0920, grad\_fn=<MseLossBackward0>)  
tensor(2381.4475, grad\_fn=<MseLossBackward0>)  
tensor(2380.8083, grad\_fn=<MseLossBackward0>)  
tensor(2380.1721, grad\_fn=<MseLossBackward0>)  
tensor(2379.5393, grad\_fn=<MseLossBackward0>)  
tensor(2378.9099, grad\_fn=<MseLossBackward0>)  
tensor(2378.2839, grad\_fn=<MseLossBackward0>)  
tensor(2377.6609, grad\_fn=<MseLossBackward0>)  
tensor(2377.0415, grad\_fn=<MseLossBackward0>)  
tensor(2376.4268, grad\_fn=<MseLossBackward0>)  
tensor(2375.8213, grad\_fn=<MseLossBackward0>)  
tensor(2375.2188, grad\_fn=<MseLossBackward0>)  
tensor(2374.6155, grad\_fn=<MseLossBackward0>)  
tensor(2373.7542, grad\_fn=<MseLossBackward0>)  
tensor(2373.1123, grad\_fn=<MseLossBackward0>)  
tensor(2372.4719, grad\_fn=<MseLossBackward0>)  
tensor(2371.8337, grad\_fn=<MseLossBackward0>)  
tensor(2371.1987, grad\_fn=<MseLossBackward0>)  
tensor(2370.5691, grad\_fn=<MseLossBackward0>)  
tensor(2369.9434, grad\_fn=<MseLossBackward0>)  
tensor(2369.3379, grad\_fn=<MseLossBackward0>)  
tensor(2368.7354, grad\_fn=<MseLossBackward0>)  
tensor(2368.1357, grad\_fn=<MseLossBackward0>)  
tensor(2367.5376, grad\_fn=<MseLossBackward0>)  
tensor(2366.9377, grad\_fn=<MseLossBackward0>)  
tensor(2366.3398, grad\_fn=<MseLossBackward0>)  
tensor(2365.7446, grad\_fn=<MseLossBackward0>)  
tensor(2364.9121, grad\_fn=<MseLossBackward0>)  
tensor(2364.2815, grad\_fn=<MseLossBackward0>)  
tensor(2363.6543, grad\_fn=<MseLossBackward0>)  
tensor(2363.0310, grad\_fn=<MseLossBackward0>)  
tensor(2362.4106, grad\_fn=<MseLossBackward0>)  
tensor(2361.7935, grad\_fn=<MseLossBackward0>)  
tensor(2361.1787, grad\_fn=<MseLossBackward0>)  
tensor(2360.5669, grad\_fn=<MseLossBackward0>)  
tensor(2359.9563, grad\_fn=<MseLossBackward0>)  
tensor(2359.3535, grad\_fn=<MseLossBackward0>)  
tensor(2358.7549, grad\_fn=<MseLossBackward0>)  
tensor(2358.1594, grad\_fn=<MseLossBackward0>)  
tensor(2357.5669, grad\_fn=<MseLossBackward0>)

tensor(2356.9775, grad\_fn=<MseLossBackward0>)  
tensor(2356.1841, grad\_fn=<MseLossBackward0>)  
tensor(2355.5654, grad\_fn=<MseLossBackward0>)  
tensor(2354.9495, grad\_fn=<MseLossBackward0>)  
tensor(2354.3408, grad\_fn=<MseLossBackward0>)  
tensor(2353.7346, grad\_fn=<MseLossBackward0>)  
tensor(2353.1316, grad\_fn=<MseLossBackward0>)  
tensor(2352.5300, grad\_fn=<MseLossBackward0>)  
tensor(2351.9309, grad\_fn=<MseLossBackward0>)  
tensor(2351.3394, grad\_fn=<MseLossBackward0>)  
tensor(2350.7532, grad\_fn=<MseLossBackward0>)  
tensor(2350.1694, grad\_fn=<MseLossBackward0>)  
tensor(2349.5889, grad\_fn=<MseLossBackward0>)  
tensor(2349.0115, grad\_fn=<MseLossBackward0>)  
tensor(2348.4370, grad\_fn=<MseLossBackward0>)  
tensor(2347.6797, grad\_fn=<MseLossBackward0>)  
tensor(2347.0759, grad\_fn=<MseLossBackward0>)  
tensor(2346.4822, grad\_fn=<MseLossBackward0>)  
tensor(2345.8916, grad\_fn=<MseLossBackward0>)  
tensor(2345.3071, grad\_fn=<MseLossBackward0>)  
tensor(2344.7175, grad\_fn=<MseLossBackward0>)  
tensor(2344.1318, grad\_fn=<MseLossBackward0>)  
tensor(2343.5520, grad\_fn=<MseLossBackward0>)  
tensor(2342.9741, grad\_fn=<MseLossBackward0>)  
tensor(2342.3982, grad\_fn=<MseLossBackward0>)  
tensor(2341.8259, grad\_fn=<MseLossBackward0>)  
tensor(2341.2549, grad\_fn=<MseLossBackward0>)  
tensor(2340.6863, grad\_fn=<MseLossBackward0>)  
tensor(2340.1182, grad\_fn=<MseLossBackward0>)  
tensor(2339.3713, grad\_fn=<MseLossBackward0>)  
tensor(2338.7761, grad\_fn=<MseLossBackward0>)  
tensor(2338.1802, grad\_fn=<MseLossBackward0>)  
tensor(2337.5908, grad\_fn=<MseLossBackward0>)  
tensor(2337.0042, grad\_fn=<MseLossBackward0>)  
tensor(2336.4197, grad\_fn=<MseLossBackward0>)  
tensor(2335.8406, grad\_fn=<MseLossBackward0>)  
tensor(2335.2729, grad\_fn=<MseLossBackward0>)  
tensor(2334.7048, grad\_fn=<MseLossBackward0>)  
tensor(2334.1355, grad\_fn=<MseLossBackward0>)  
tensor(2333.5708, grad\_fn=<MseLossBackward0>)  
tensor(2333.0073, grad\_fn=<MseLossBackward0>)

```
tensor(2332.4473, grad_fn=<MseLossBackward0>)  
tensor(2331.8894, grad_fn=<MseLossBackward0>)  
tensor(2331.1370, grad_fn=<MseLossBackward0>)  
tensor(2330.5532, grad_fn=<MseLossBackward0>)  
tensor(2329.9731, grad_fn=<MseLossBackward0>)  
tensor(2329.3979, grad_fn=<MseLossBackward0>)  
tensor(2328.8264, grad_fn=<MseLossBackward0>)  
tensor(2328.2559, grad_fn=<MseLossBackward0>)  
tensor(2327.6873, grad_fn=<MseLossBackward0>)  
tensor(2327.1230, grad_fn=<MseLossBackward0>)  
tensor(2326.5615, grad_fn=<MseLossBackward0>)  
tensor(2326.0034, grad_fn=<MseLossBackward0>)  
tensor(2325.4482, grad_fn=<MseLossBackward0>)  
tensor(2324.8960, grad_fn=<MseLossBackward0>)  
tensor(2324.3457, grad_fn=<MseLossBackward0>)  
tensor(2323.8000, grad_fn=<MseLossBackward0>)  
tensor(2323.0796, grad_fn=<MseLossBackward0>)  
tensor(2322.5195, grad_fn=<MseLossBackward0>)  
tensor(2321.9587, grad_fn=<MseLossBackward0>)  
tensor(2321.3984, grad_fn=<MseLossBackward0>)  
tensor(2320.8430, grad_fn=<MseLossBackward0>)  
tensor(2320.2932, grad_fn=<MseLossBackward0>)  
tensor(2319.7451, grad_fn=<MseLossBackward0>)  
tensor(2319.2026, grad_fn=<MseLossBackward0>)  
tensor(2318.6626, grad_fn=<MseLossBackward0>)  
tensor(2318.1272, grad_fn=<MseLossBackward0>)  
tensor(2317.5950, grad_fn=<MseLossBackward0>)  
tensor(2317.0642, grad_fn=<MseLossBackward0>)  
tensor(2316.5383, grad_fn=<MseLossBackward0>)  
tensor(2316.0144, grad_fn=<MseLossBackward0>)  
tensor(2315.3186, grad_fn=<MseLossBackward0>)  
tensor(2314.7651, grad_fn=<MseLossBackward0>)  
tensor(2314.2065, grad_fn=<MseLossBackward0>)  
tensor(2313.6548, grad_fn=<MseLossBackward0>)  
tensor(2313.1030, grad_fn=<MseLossBackward0>)  
tensor(2312.5525, grad_fn=<MseLossBackward0>)  
tensor(2312.0039, grad_fn=<MseLossBackward0>)  
tensor(2311.4597, grad_fn=<MseLossBackward0>)  
tensor(2310.9177, grad_fn=<MseLossBackward0>)  
tensor(2310.3774, grad_fn=<MseLossBackward0>)  
tensor(2309.8406, grad_fn=<MseLossBackward0>)
```



```

tensor(2309.3066, grad_fn=<MseLossBackward0>)
tensor(2308.7769, grad_fn=<MseLossBackward0>)
tensor(2308.2498, grad_fn=<MseLossBackward0>)
tensor(2307.5374, grad_fn=<MseLossBackward0>)
tensor(2306.9841, grad_fn=<MseLossBackward0>)
tensor(2306.4329, grad_fn=<MseLossBackward0>)
tensor(2305.8892, grad_fn=<MseLossBackward0>)
tensor(2305.3472, grad_fn=<MseLossBackward0>)
tensor(2304.8159, grad_fn=<MseLossBackward0>)
tensor(2304.2869, grad_fn=<MseLossBackward0>)
tensor(2303.7603, grad_fn=<MseLossBackward0>)
tensor(2303.2358, grad_fn=<MseLossBackward0>)
tensor(2302.7139, grad_fn=<MseLossBackward0>)
tensor(2302.1946, grad_fn=<MseLossBackward0>)
tensor(2301.6787, grad_fn=<MseLossBackward0>)
tensor(2301.1660, grad_fn=<MseLossBackward0>)
tensor(2300.6550, grad_fn=<MseLossBackward0>)

```

```

In [18]: y_pred, _, _=model(Xte)
         #performance metric for classification
         #print(f'ACC:{accuracy_score(yte.detach().numpy(),y_pred.detach().numpy())>0.5}}') #classification
         #performance metric for regression
         print(f'MSE:{mean_squared_error(yte.detach().numpy(),y_pred.detach().numpy())}') #regression

```

MSE: 2391.4755859375

In comparison with TSK model previous used, it was possible to slightly improve the results, going from a MSE of 2476.79 to 2391.48. The change was not significative. Given the range of the target values, the MSE obtain represents a large value, not giving confidence for a certain predicted value. When visualizing the error (as plotted in the chart bellow) it is possible to notice that, bigger the value of the targer, beter it predicts it. It is possible to verify that the predicted values tend to map the trend (if the real value is bigger, the predicted one tend to be bigger as well). The problem is that it seems to be affected by a large "noise", responsible for the large MSE value obtained.

```

In [19]: # Plot predictions vs actual
         # Converter y_pred para numpy e flatten
         y_pred_np = y_pred.detach().numpy().flatten()
         yte_np = yte.detach().numpy()

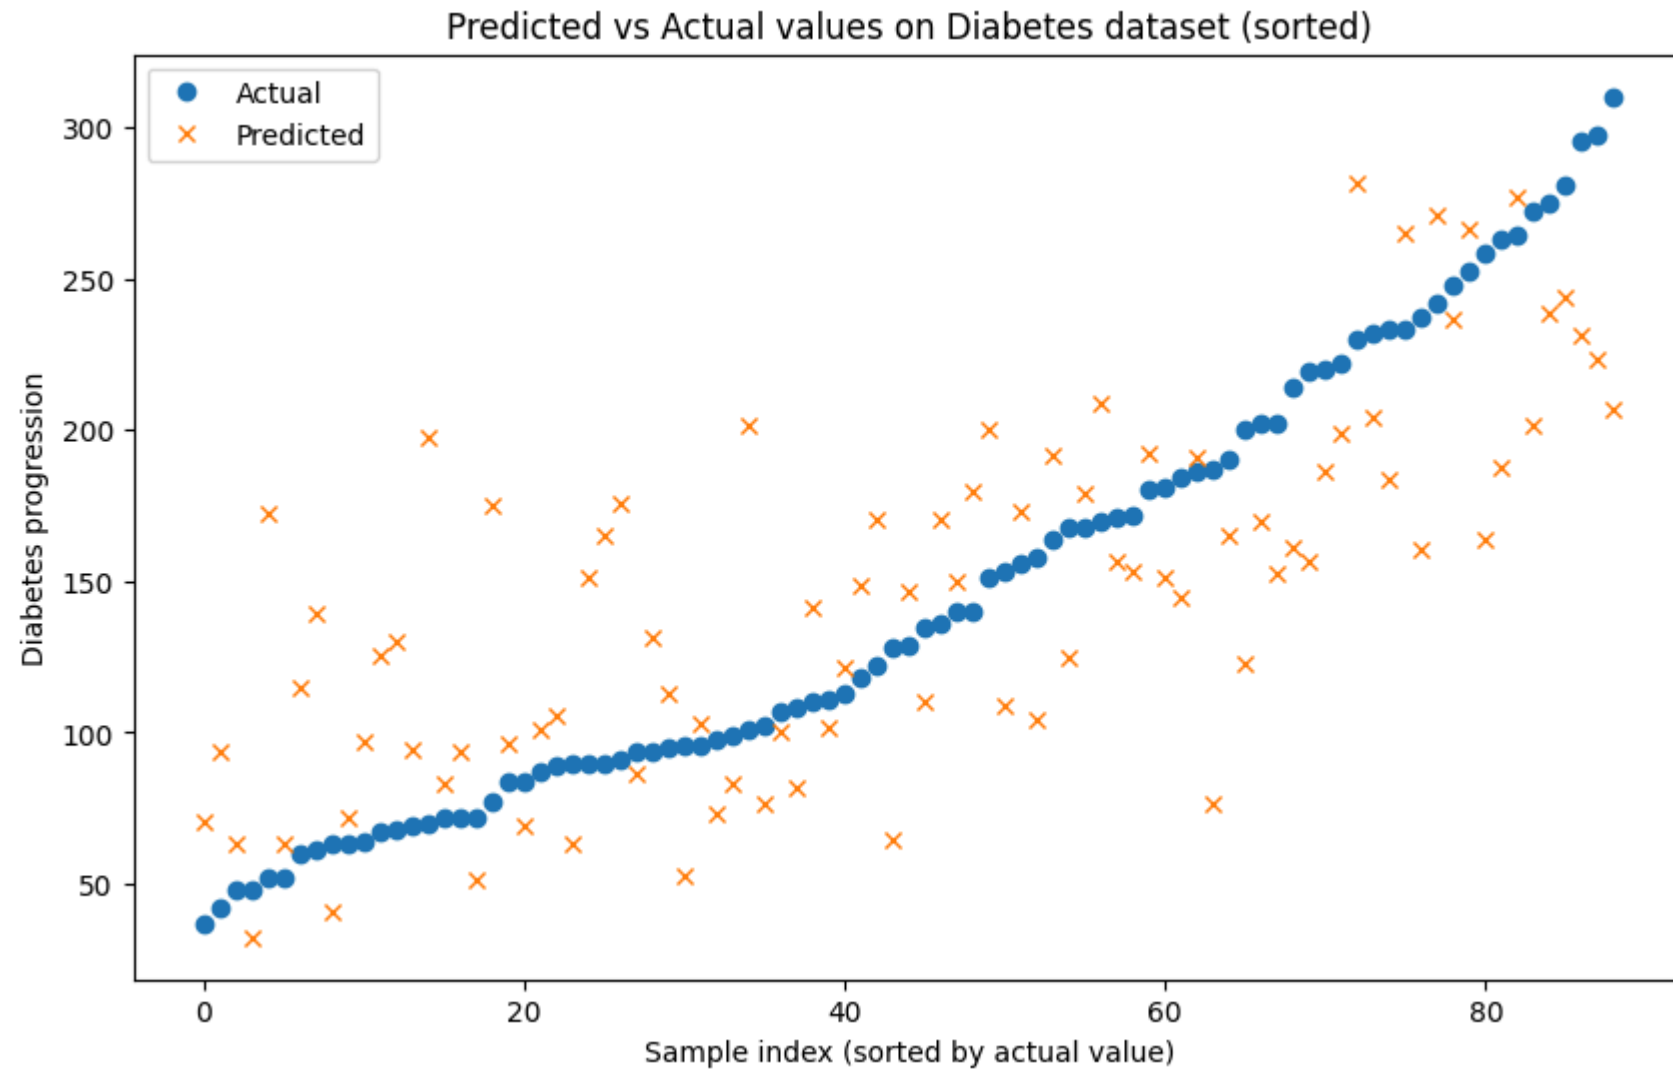
         # Obter índices que ordenam yte
         sort_idx = np.argsort(yte_np)

```

```
# Ordenar yte e y_pred segundo esses índices
yte_sorted = yte_np[sort_idx]
y_pred_sorted = y_pred_np[sort_idx]

# Plot
plt.figure(figsize=(10,6))
plt.plot(range(len(yte_sorted)), yte_sorted, label="Actual", marker="o", linestyle='')
plt.plot(range(len(y_pred_sorted)), y_pred_sorted, label="Predicted", marker="x", linestyle='')

plt.xlabel("Sample index (sorted by actual value)")
plt.ylabel("Diabetes progression")
plt.title("Predicted vs Actual values on Diabetes dataset (sorted)")
plt.legend()
plt.show()
```



```
In [1]: import numpy as np
        from sklearn import datasets
        from sklearn.preprocessing import StandardScaler
        from sklearn.model_selection import train_test_split
        from sklearn.metrics import mean_squared_error, accuracy_score, classification_report
        import matplotlib.pyplot as plt
        import torch.nn.functional as F
        import torch
        import torch.nn as nn
        import torch.optim as optim
        from torch.utils.data import TensorDataset, DataLoader
        import pandas
```

```
In [2]: # CHOOSE DATASET

        # Binary classification dataset
        diabetes = datasets.load_diabetes(as_frame=True)

        # Regression dataset
        #data = datasets.fetch_openml(name="boston", version=1, as_frame=True)

        X = diabetes.data.values
        y = diabetes.target.values

        print("Shape:", X.shape)

        print(diabetes.data.head(), "\n \n") # first rows of features
        print(diabetes.target.head()) # first rows of target
```

Shape: (442, 10)

	age	sex	bmi	bp	s1	s2	s3	\
0	0.038076	0.050680	0.061696	0.021872	-0.044223	-0.034821	-0.043401	
1	-0.001882	-0.044642	-0.051474	-0.026328	-0.008449	-0.019163	0.074412	
2	0.085299	0.050680	0.044451	-0.005670	-0.045599	-0.034194	-0.032356	
3	-0.089063	-0.044642	-0.011595	-0.036656	0.012191	0.024991	-0.036038	
4	0.005383	-0.044642	-0.036385	0.021872	0.003935	0.015596	0.008142	

	s4	s5	s6
0	-0.002592	0.019907	-0.017646
1	-0.039493	-0.068332	-0.092204
2	-0.002592	0.002861	-0.025930
3	0.034309	0.022688	-0.009362
4	-0.002592	-0.031988	-0.046641

0	151.0
1	75.0
2	141.0
3	206.0
4	135.0

Name: target, dtype: float64

```
In [3]: #train test splitting
test_size=0.2
Xtr, Xte, ytr, yte = train_test_split(X, y, test_size=test_size, random_state=42)
```

```
In [4]: # Standardize features
scaler=StandardScaler()
Xtr= scaler.fit_transform(Xtr)
Xte= scaler.transform(Xte)
```

A fixed seed was added to this code to ensure the reproducibility of the analysis. This allowed for manual tuning of the hyperparameters to achieve better model training.

```
In [5]: import random

seed = 42
torch.manual_seed(seed)
```

```
np.random.seed(seed)
random.seed(seed)

# Para GPU
torch.cuda.manual_seed(seed)
torch.cuda.manual_seed_all(seed)

# Tornar CUDA determinístico
torch.backends.cudnn.deterministic = True
torch.backends.cudnn.benchmark = False
```

```
In [6]: class MLP(nn.Module):
        def __init__(self, input_size, output_size=1, dropout_prob=0.5):
            super(MLP, self).__init__()

            self.fc1 = nn.Linear(input_size, 64)
            self.fc2 = nn.Linear(64, 64)
            self.fc3 = nn.Linear(64, 64)
            self.fc4 = nn.Linear(64, 64)
            self.out = nn.Linear(64, output_size)

            self.dropout = nn.Dropout(p=dropout_prob)

        def forward(self, x):
            x = F.relu(self.fc1(x))
            x = self.dropout(x)

            x = F.relu(self.fc2(x))
            x = self.dropout(x)

            x = F.relu(self.fc3(x))
            x = self.dropout(x)

            x = F.relu(self.fc4(x))
            x = self.dropout(x)

            x = self.out(x)
            return x
```

```
In [7]: num_epochs=190
        lr=0.01
        dropout=0.25
        batch_size=128
```

This model was trained on the GPU and then transferred to the CPU for use with NumPy. Given the low number of parameters (i.e., the model's low complexity), the time required on the CPU was similar to that on the GPU.

```
In [8]: # Model, Loss, Optimizer
        device = torch.device("cuda" if torch.cuda.is_available() else "cpu")
        #device = "cpu" # force to use CPU
        print(device)
```

cuda

```
In [9]: Xtr = torch.tensor(Xtr, dtype=torch.float32).to(device)
        ytr = torch.tensor(ytr, dtype=torch.float32).to(device)
        Xte = torch.tensor(Xte, dtype=torch.float32).to(device)
        yte = torch.tensor(yte, dtype=torch.float32).to("cpu")

        # Wrap Xtr and ytr into a dataset
        train_dataset = TensorDataset(Xtr, ytr)

        # Create DataLoader
        train_dataloader = DataLoader(train_dataset, batch_size=batch_size, shuffle=True)
```

```
In [10]: model = MLP(input_size=Xtr.shape[1], dropout_prob=dropout).to(device)
        criterion = nn.BCEWithLogitsLoss() # for binary classification
        criterion = nn.MSELoss() #for regression
        optimizer = optim.Adam(model.parameters(), lr=lr)
```

The model was implemented as a fully connected neural network (MLP) with four hidden layers of 64 neurons each, using ReLU activation functions. The network takes the input features of the dataset and outputs a single value for regression (Diabetes Progression).

```
In [11]: # Training Loop
        import time
        start_time = time.time()
        for epoch in range(num_epochs):
```

```
model.train()
epoch_loss = 0.0

for batch_x, batch_y in train_dataloader:
    batch_x = batch_x.to(device)
    batch_y = batch_y.to(device)

    logits = model(batch_x)
    loss = criterion(logits, batch_y.view(-1, 1))

    optimizer.zero_grad()
    loss.backward()
    optimizer.step()

    epoch_loss += loss.item()

avg_loss = epoch_loss / len(train_dataloader)
print(f"Epoch [{epoch+1}/{num_epochs}], Loss: {avg_loss:.4f}")

end_time = time.time()
print(f"Training time: {end_time - start_time:.2f} seconds")
```



Epoch [1/190], Loss: 29648.7161  
Epoch [2/190], Loss: 27738.2480  
Epoch [3/190], Loss: 19526.0052  
Epoch [4/190], Loss: 10916.4779  
Epoch [5/190], Loss: 7908.8825  
Epoch [6/190], Loss: 6315.4227  
Epoch [7/190], Loss: 6881.3011  
Epoch [8/190], Loss: 5394.8262  
Epoch [9/190], Loss: 5833.4985  
Epoch [10/190], Loss: 4486.8579  
Epoch [11/190], Loss: 4801.0705  
Epoch [12/190], Loss: 4380.2632  
Epoch [13/190], Loss: 4705.3464  
Epoch [14/190], Loss: 4574.8371  
Epoch [15/190], Loss: 4115.9676  
Epoch [16/190], Loss: 4263.3736  
Epoch [17/190], Loss: 4220.0199  
Epoch [18/190], Loss: 3679.3091  
Epoch [19/190], Loss: 3745.2568  
Epoch [20/190], Loss: 3896.5340  
Epoch [21/190], Loss: 4014.5900  
Epoch [22/190], Loss: 3893.0127  
Epoch [23/190], Loss: 3573.5227  
Epoch [24/190], Loss: 3891.4432  
Epoch [25/190], Loss: 3775.5807  
Epoch [26/190], Loss: 3753.8115  
Epoch [27/190], Loss: 3697.5588  
Epoch [28/190], Loss: 3638.7832  
Epoch [29/190], Loss: 4002.4415  
Epoch [30/190], Loss: 3921.0071  
Epoch [31/190], Loss: 3542.5956  
Epoch [32/190], Loss: 3535.1024  
Epoch [33/190], Loss: 3598.4771  
Epoch [34/190], Loss: 3983.6483  
Epoch [35/190], Loss: 3522.1357  
Epoch [36/190], Loss: 3819.0868  
Epoch [37/190], Loss: 3664.6637  
Epoch [38/190], Loss: 3815.6850  
Epoch [39/190], Loss: 3901.3689  
Epoch [40/190], Loss: 3601.0779  
Epoch [41/190], Loss: 3759.6592

Epoch [42/190], Loss: 3391.2233  
Epoch [43/190], Loss: 3242.0986  
Epoch [44/190], Loss: 3577.6099  
Epoch [45/190], Loss: 3480.0977  
Epoch [46/190], Loss: 3555.8261  
Epoch [47/190], Loss: 3476.3621  
Epoch [48/190], Loss: 3268.5854  
Epoch [49/190], Loss: 3618.4499  
Epoch [50/190], Loss: 3584.3310  
Epoch [51/190], Loss: 3675.5367  
Epoch [52/190], Loss: 3722.2685  
Epoch [53/190], Loss: 3227.3150  
Epoch [54/190], Loss: 2974.1118  
Epoch [55/190], Loss: 3562.6582  
Epoch [56/190], Loss: 3198.2245  
Epoch [57/190], Loss: 3238.8750  
Epoch [58/190], Loss: 3376.6782  
Epoch [59/190], Loss: 3346.6515  
Epoch [60/190], Loss: 3468.1538  
Epoch [61/190], Loss: 3152.5792  
Epoch [62/190], Loss: 3484.5428  
Epoch [63/190], Loss: 3397.3236  
Epoch [64/190], Loss: 3454.3319  
Epoch [65/190], Loss: 3955.6242  
Epoch [66/190], Loss: 3623.3734  
Epoch [67/190], Loss: 3310.0218  
Epoch [68/190], Loss: 3456.5499  
Epoch [69/190], Loss: 3696.4705  
Epoch [70/190], Loss: 3757.5662  
Epoch [71/190], Loss: 3274.0444  
Epoch [72/190], Loss: 3489.4048  
Epoch [73/190], Loss: 3018.0824  
Epoch [74/190], Loss: 3315.5653  
Epoch [75/190], Loss: 3203.5889  
Epoch [76/190], Loss: 3338.0278  
Epoch [77/190], Loss: 3350.0305  
Epoch [78/190], Loss: 3378.0791  
Epoch [79/190], Loss: 3423.2143  
Epoch [80/190], Loss: 3527.6632  
Epoch [81/190], Loss: 3268.2509  
Epoch [82/190], Loss: 3227.4768

Epoch [83/190], Loss: 3589.3934  
Epoch [84/190], Loss: 3299.7594  
Epoch [85/190], Loss: 3385.5459  
Epoch [86/190], Loss: 3498.3327  
Epoch [87/190], Loss: 3477.1771  
Epoch [88/190], Loss: 3174.3293  
Epoch [89/190], Loss: 3714.6216  
Epoch [90/190], Loss: 3343.3033  
Epoch [91/190], Loss: 3414.0710  
Epoch [92/190], Loss: 3248.5864  
Epoch [93/190], Loss: 3488.2712  
Epoch [94/190], Loss: 3373.2797  
Epoch [95/190], Loss: 3372.9124  
Epoch [96/190], Loss: 3554.2501  
Epoch [97/190], Loss: 3368.5915  
Epoch [98/190], Loss: 3215.7687  
Epoch [99/190], Loss: 3260.8147  
Epoch [100/190], Loss: 3192.8174  
Epoch [101/190], Loss: 3423.2056  
Epoch [102/190], Loss: 2991.1534  
Epoch [103/190], Loss: 3152.7769  
Epoch [104/190], Loss: 3303.9213  
Epoch [105/190], Loss: 3282.2486  
Epoch [106/190], Loss: 3331.6805  
Epoch [107/190], Loss: 3279.8656  
Epoch [108/190], Loss: 3262.2849  
Epoch [109/190], Loss: 3125.2458  
Epoch [110/190], Loss: 3407.5113  
Epoch [111/190], Loss: 3498.1632  
Epoch [112/190], Loss: 3553.5989  
Epoch [113/190], Loss: 3089.8211  
Epoch [114/190], Loss: 3075.9001  
Epoch [115/190], Loss: 3539.4534  
Epoch [116/190], Loss: 2961.9173  
Epoch [117/190], Loss: 3655.1819  
Epoch [118/190], Loss: 3049.0596  
Epoch [119/190], Loss: 3615.5492  
Epoch [120/190], Loss: 3256.5053  
Epoch [121/190], Loss: 3308.7235  
Epoch [122/190], Loss: 3240.1512  
Epoch [123/190], Loss: 3193.3521

Epoch [124/190], Loss: 3216.8752  
Epoch [125/190], Loss: 3073.2205  
Epoch [126/190], Loss: 3247.6636  
Epoch [127/190], Loss: 3242.5971  
Epoch [128/190], Loss: 3164.6580  
Epoch [129/190], Loss: 3109.4061  
Epoch [130/190], Loss: 3200.9823  
Epoch [131/190], Loss: 2903.9763  
Epoch [132/190], Loss: 3002.0817  
Epoch [133/190], Loss: 3106.7635  
Epoch [134/190], Loss: 3022.0481  
Epoch [135/190], Loss: 2955.9606  
Epoch [136/190], Loss: 2726.1379  
Epoch [137/190], Loss: 2718.2030  
Epoch [138/190], Loss: 2962.2109  
Epoch [139/190], Loss: 2859.1110  
Epoch [140/190], Loss: 2965.7663  
Epoch [141/190], Loss: 2804.8240  
Epoch [142/190], Loss: 3106.8247  
Epoch [143/190], Loss: 2979.2811  
Epoch [144/190], Loss: 3120.4154  
Epoch [145/190], Loss: 3010.3355  
Epoch [146/190], Loss: 2924.7756  
Epoch [147/190], Loss: 3123.8313  
Epoch [148/190], Loss: 3109.5750  
Epoch [149/190], Loss: 3098.7642  
Epoch [150/190], Loss: 2857.3700  
Epoch [151/190], Loss: 2996.3384  
Epoch [152/190], Loss: 3129.6170  
Epoch [153/190], Loss: 2759.2476  
Epoch [154/190], Loss: 3000.5859  
Epoch [155/190], Loss: 2925.3764  
Epoch [156/190], Loss: 2861.1976  
Epoch [157/190], Loss: 2986.6821  
Epoch [158/190], Loss: 2743.7799  
Epoch [159/190], Loss: 3126.5764  
Epoch [160/190], Loss: 2914.5562  
Epoch [161/190], Loss: 2880.0532  
Epoch [162/190], Loss: 2857.0591  
Epoch [163/190], Loss: 2817.7529  
Epoch [164/190], Loss: 2701.1991

```
Epoch [165/190], Loss: 3057.9538
Epoch [166/190], Loss: 2827.3732
Epoch [167/190], Loss: 2802.8832
Epoch [168/190], Loss: 3131.7845
Epoch [169/190], Loss: 3075.6080
Epoch [170/190], Loss: 2931.3674
Epoch [171/190], Loss: 2838.9221
Epoch [172/190], Loss: 3041.7693
Epoch [173/190], Loss: 3094.6465
Epoch [174/190], Loss: 3209.8215
Epoch [175/190], Loss: 3203.1895
Epoch [176/190], Loss: 3175.8590
Epoch [177/190], Loss: 2773.8187
Epoch [178/190], Loss: 3101.5061
Epoch [179/190], Loss: 3159.9993
Epoch [180/190], Loss: 2911.1960
Epoch [181/190], Loss: 2810.7509
Epoch [182/190], Loss: 2745.5351
Epoch [183/190], Loss: 2873.0893
Epoch [184/190], Loss: 3086.1829
Epoch [185/190], Loss: 3048.7624
Epoch [186/190], Loss: 2893.8034
Epoch [187/190], Loss: 2915.0261
Epoch [188/190], Loss: 2572.4602
Epoch [189/190], Loss: 2886.5369
Epoch [190/190], Loss: 2821.7037
Training time: 1.99 seconds
```

```
In [12]: y_pred=model(Xte).cpu() # só nesta altura volta ao CPU
          #print(f'ACC:{accuracy_score(yte.detach().numpy(),y_pred.detach().numpy())>0.5}') #classification

          print(f'MSE:{mean_squared_error(yte.detach().numpy(),y_pred.detach().numpy())}') #regression
```

MSE: 2814.833251953125

After tuning the hyperparameters, the best configuration was: num\_epochs = 200, lr = 0.02, dropout = 0.2, and batch\_size = 128. A dropout with a probability of 0.2 was applied after each layer to mitigate overfitting. In order to validate the model the MSE was computed. The model was evaluated using the Mean Squared Error (MSE), which resulted in a high value of 2815. This indicates that, despite the training and parameter optimization, the current approach is insufficient to accurately predict Diabetes Progression. The high error may be attributed to

the limited size of the dataset, the low number of features. In comparison with previous approaches, such as ANFIS or TSK models, this approach exhibited the worst performance.


```
In [13]: # Plot predictions vs actual
# Converter y_pred para numpy e flatten
y_pred_np = y_pred.detach().numpy().flatten()
yte_np = yte.detach().numpy()

# Obter índices que ordenam yte
sort_idx = np.argsort(yte_np)

# Ordenar yte e y_pred segundo esses índices
yte_sorted = yte_np[sort_idx]
y_pred_sorted = y_pred_np[sort_idx]

# Plot
plt.figure(figsize=(10,6))
plt.plot(range(len(yte_sorted)), yte_sorted, label="Actual", marker="o", linestyle='')
plt.plot(range(len(y_pred_sorted)), y_pred_sorted, label="Predicted", marker="x", linestyle='')

plt.xlabel("Sample index (sorted by actual value)")
plt.ylabel("Diabetes progression")
plt.title("Predicted vs Actual values on Diabetes dataset (sorted)")
plt.legend()
plt.show()
```



```
In [96]: import numpy as np
from sklearn.datasets import fetch_openml
from sklearn.preprocessing import StandardScaler
from sklearn.model_selection import train_test_split
from sklearn.metrics import mean_squared_error, accuracy_score, classification_report, confusion_matrix

import skfuzzy as fuzz
import matplotlib.pyplot as plt
import torch
import torch.nn as nn
import torch.optim as optim
import pandas
```

Importation of Dataset 2

```
In [97]: # CHOOSE DATASET

# Binary classification dataset
diabetes = fetch_openml("diabetes", version = 1, as_frame=True)

X = diabetes.data.values
y = diabetes.target.values

y = np.where(y == "tested_positive", 1, 0)

print("Shape:", X.shape)

print(diabetes.data.head(), "\n \n") # first rows of features
print(diabetes.target.head()) # first rows of target
```



Shape: (768, 8)

	preg	plas	pres	skin	insu	mass	pedi	age
0	6	148	72	35	0	33.6	0.627	50
1	1	85	66	29	0	26.6	0.351	31
2	8	183	64	0	0	23.3	0.672	32
3	1	89	66	23	94	28.1	0.167	21
4	0	137	40	35	168	43.1	2.288	33

0 tested\_positive

1 tested\_negative

2 tested\_positive

3 tested\_negative

4 tested\_positive

Name: class, dtype: category

Categories (2, object): ['tested\_negative', 'tested\_positive']

```
In [98]: #train test splitting
test_size=0.2
Xtr, Xte, ytr, yte = train_test_split(X, y, test_size=test_size, random_state=42)
```

```
In [99]: # Standardize features
scaler=StandardScaler()
Xtr= scaler.fit_transform(Xtr)
Xte= scaler.transform(Xte)
```

In order to be able to compare the results, the number of clusters and value of m used, was the same as the one in the previous assigment ( n\_clusters = 4; m=1.1)

```
In [100... # Number of clusters
n_clusters = 4
m=1.1

# Concatenate target for clustering
Xexp=np.concatenate([Xtr, ytr.reshape(-1, 1)], axis=1)
#Xexp=Xtr

# Transpose data for skfuzzy (expects features x samples)
Xexp_T = Xexp.T
```

```
# Fuzzy C-means clustering
centers, u, u0, d, jm, p, fpc = fuzz.cluster.cmeans(
    Xexp_T, n_clusters, m=m, error=0.005, maxiter=1000, init=None,
)
```

In [101... centers.shape

Out[101... (4, 9)

```
# Compute sigma (spread) for each cluster
sigmas = []
for j in range(n_clusters):
    # membership weights for cluster j, raised to m
    u_j = u[j, :] ** m
    # weighted variance for each feature
    var_j = np.average((Xexp - centers[j])**2, axis=0, weights=u_j)
    sigma_j = np.sqrt(var_j)
    sigmas.append(sigma_j)
sigmas=np.array(sigmas)
```


```
# Hard clustering from fuzzy membership
cluster_labels = np.argmax(u, axis=0)
print("Fuzzy partition coefficient (FPC):", fpc)

# Plot first two features with fuzzy membership
plt.figure(figsize=(8,6))
for j in range(n_clusters):
    plt.scatter(
        Xexp[cluster_labels == j, 0],          # Feature 1
        Xexp[cluster_labels == j, 1],          # Feature 2
        alpha=u[j, :],                          # transparency ~ membership
        label=f'Cluster {j}'
    )

plt.title("Fuzzy C-Means Clustering (with membership degree)")
plt.xlabel("Feature 1")
plt.ylabel("Feature 2")
```


```
plt.legend()
plt.show()
```

Fuzzy partition coefficient (FPC): 0.9091786112603655



```
In [104... # Plot first two features with cluster assignments
plt.figure(figsize=(8,6))
for j in range(n_clusters):
    plt.scatter(
```

```
        Xexp[cluster_labels == j, 0],  
        Xexp[cluster_labels == j, 1],  
        label=f'Cluster {j}'  
    )  
  
plt.title("Fuzzy C-Means Clustering (CRISPEN)")  
plt.xlabel("Feature 1")  
plt.ylabel("Feature 2")  
plt.legend()  
plt.show()
```




```
In [105... # Gaussian formula
def gaussian(x, mu, sigma):
    return np.exp(-0.5 * ((x - mu)/sigma)**2)

lin=np.linspace(-2, 4, 500)
plt.figure(figsize=(8,6))
```

```
y_aux=[]
feature=0
for j in range(n_clusters):
    # Compute curves
    y_aux.append(gaussian(lin, centers[j,feature], sigmas[j,feature]))

    # Plot
    plt.plot(lin, y_aux[j], label=f"Gaussian  $\mu$ ={np.round(centers[j,feature],2)},  $\sigma$ ={np.round(sigmas[j,feature],2)}")

plt.title("Projection of the membership functions on Feature 2")
plt.xlabel("Feature 1")
plt.ylabel("Degree of Membership")
plt.legend()
plt.grid(True)
plt.show()
```



In [106...

```
# -----
# Gaussian Membership Function
# -----
class GaussianMF(nn.Module):
    def __init__(self, centers, sigmas, agg_prob):
        super().__init__()
        self.centers = nn.Parameter(torch.tensor(centers, dtype=torch.float32))
```

```

self.sigmas = nn.Parameter(torch.tensor(sigmas, dtype=torch.float32))
self.agg_prob=agg_prob

def forward(self, x):
    # Expand for broadcasting
    # x: (batch, 1, n_dims), centers: (1, n_rules, n_dims), sigmas: (1, n_rules, n_dims)
    diff = abs((x.unsqueeze(1) - self.centers.unsqueeze(0))/self.sigmas.unsqueeze(0)) #(batch, n_rules, n_dims)

    # Aggregation
    if self.agg_prob:
        dist = torch.norm(diff, dim=-1) # (batch, n_rules) # probabilistic intersection
    else:
        dist = torch.max(diff, dim=-1).values # (batch, n_rules) # min intersection (min instersection of normal funtion

    return torch.exp(-0.5 * dist ** 2)

# -----
# TSK Model
# -----
class TSK(nn.Module):
    def __init__(self, n_inputs, n_rules, centers, sigmas,agg_prob=False):
        super().__init__()
        self.n_inputs = n_inputs
        self.n_rules = n_rules

        # Antecedents (Gaussian MFs)

        self.mfs=GaussianMF(centers, sigmas,agg_prob)

        # Consequents (Linear functions of inputs)
        # Each rule has coeffs for each input + bias
        self.consequents = nn.Parameter(
            torch.randn(n_inputs + 1,n_rules)
        )

    def forward(self, x):
        # x: (batch, n_inputs)
        batch_size = x.shape[0]

        # Compute membership values for each input feature

```



```

# firing_strengths: (batch, n_rules)
firing_strengths = self.mfs(x)

# Normalize memberships
# norm_fs: (batch, n_rules)
norm_fs = firing_strengths / (firing_strengths.sum(dim=1, keepdim=True) + 1e-9)

# Consequent output (linear model per rule)
x_aug = torch.cat([x, torch.ones(batch_size, 1)], dim=1) # add bias

rule_outputs = torch.einsum("br,rk->bk", x_aug, self.consequents) # (batch, rules)
# Weighted sum
output = torch.sum(norm_fs * rule_outputs, dim=1, keepdim=True)

return output, norm_fs, rule_outputs

```

In [107...

```

# -----
# Least Squares Solver for Consequents (TSK)
# -----
def train_ls(model, X, y):
    with torch.no_grad():
        _, norm_fs, _ = model(X)

        # Design matrix for LS: combine normalized firing strengths with input
        X_aug = torch.cat([X, torch.ones(X.shape[0], 1)], dim=1)

        Phi = torch.einsum("br,bi->bri", X_aug, norm_fs).reshape(X.shape[0], -1)

        # Solve LS: consequents = (Phi^T Phi)^-1 Phi^T y

        theta = torch.linalg.lstsq(Phi, y).solution

        model.consequents.data = theta.reshape(model.consequents.shape)

```

In [108...

```

# -----
# Gradient Descent Training
# -----

```

```
def train_gd(model, X, y, epochs=100, lr=1e-3):
    optimizer = optim.Adam(model.parameters(), lr=lr)
    criterion = nn.MSELoss()
    for _ in range(epochs):
        optimizer.zero_grad()
        y_pred, _, _ = model(X)
        loss = criterion(y_pred, y)
        #print(loss)
        loss.backward()
        optimizer.step()
```

```
In [109... # -----
# Hybrid Training (Classic ANFIS)
# -----
def train_hybrid_anfis(model, X, y, max_iters=10, gd_epochs=20, lr=1e-3): #10, 20, 1e-3
    train_ls(model, X, y)
    for _ in range(max_iters):
        # Step A: GD on antecedents (freeze consequents)
        model.consequents.requires_grad = False
        train_gd(model, X, y, epochs=gd_epochs, lr=lr)

        # Step B: LS on consequents (freeze antecedents)
        model.consequents.requires_grad = True
        model.mfs.requires_grad = False
        train_ls(model, X, y)

        # Re-enable antecedents
        model.mfs.requires_grad = True
```

```
In [110... # -----
# Alternative Hybrid Training (LS+ gradient descent on all)
# -----
def train_hybrid(model, X, y, epochs=100, lr=1e-5): #def; 100, 4
    # Step 1: LS for consequents
    train_ls(model, X, y)
    # Step 2: GD fine-tuning
    train_gd(model, X, y, epochs=epochs, lr=lr)
```

```
In [111... # Build model
model = TSK(n_inputs=Xtr.shape[1], n_rules=n_clusters, centers=centers[:, :-1], sigmas=sigmas[:, :-1])
```

```
Xtr = torch.tensor(Xtr, dtype=torch.float32)
ytr = torch.tensor(ytr, dtype=torch.float32)
Xte = torch.tensor(Xte, dtype=torch.float32)
yte = torch.tensor(yte, dtype=torch.float32)
```

```
In [112... # Training with LS:
#train_ls(model, Xtr, ytr.reshape(-1,1))
train_hybrid_anfis(model, Xtr, ytr.reshape(-1,1), max_iters=13, gd_epochs=11, lr=1e-4) #10 20 3
```

```
In [113... thr = 0.6 #threshold to tune

y_pred, _, _ = model(Xte)
y_true = yte.detach().numpy()
y_pred_bin = (y_pred.detach().numpy() > thr).astype(int) # binary predictions

# Accuracy
acc = accuracy_score(y_true, y_pred_bin)
print(f'ACC:{accuracy_score(y_true,y_pred_bin)}') #classification

cm = confusion_matrix(y_true, y_pred_bin)
print("Confusion Matrix:")
print(cm)

#print(f'ACC:{mean_squared_error(yte.detach().numpy(),y_pred.detach().numpy())}') #regression
```

ACC:0.8116883116883117

Confusion Matrix:

```
[[93  6]
 [23 32]]
```

After manual tuning, the best model was achieved after changing the hyperparameters to: max\_iters=13, gd\_epochs=11, lr=1e-4. In comparison with TSK model previous used, it was possible to slightly improve the results, going from a accuracy score of 79.87% to 81.17%. The change was not significative, beeing able to only reduce 2 FP values (from 8 to 6). When visualizing the error (as plotted in the chart bellow) it is possible to notice that the model mutch better classifies for negative examples, represent by a clear separation in the predicted target distribution. To predict the positive exmples to model have more dificulty, presenting a recall score of 58.19% (TP/(TP+FN)).

```
In [114... # Convert y_pred to numpy and flatten
y_pred_np = y_pred.detach().numpy().flatten()
```

```
y_pred_bin_np = y_pred_bin
yte_np = yte.detach().numpy()

# Get indices that sort yte
sort_idx = np.argsort(yte_np)

# Sort yte and y_pred
yte_sorted = yte_np[sort_idx]
y_pred_sorted = y_pred_np[sort_idx]

# Determine the index where actual target switches from 0 to 1
frontier_idx = np.argmax(yte_sorted == 1) # first occurrence of 1

# Plot
plt.figure(figsize=(10,6))

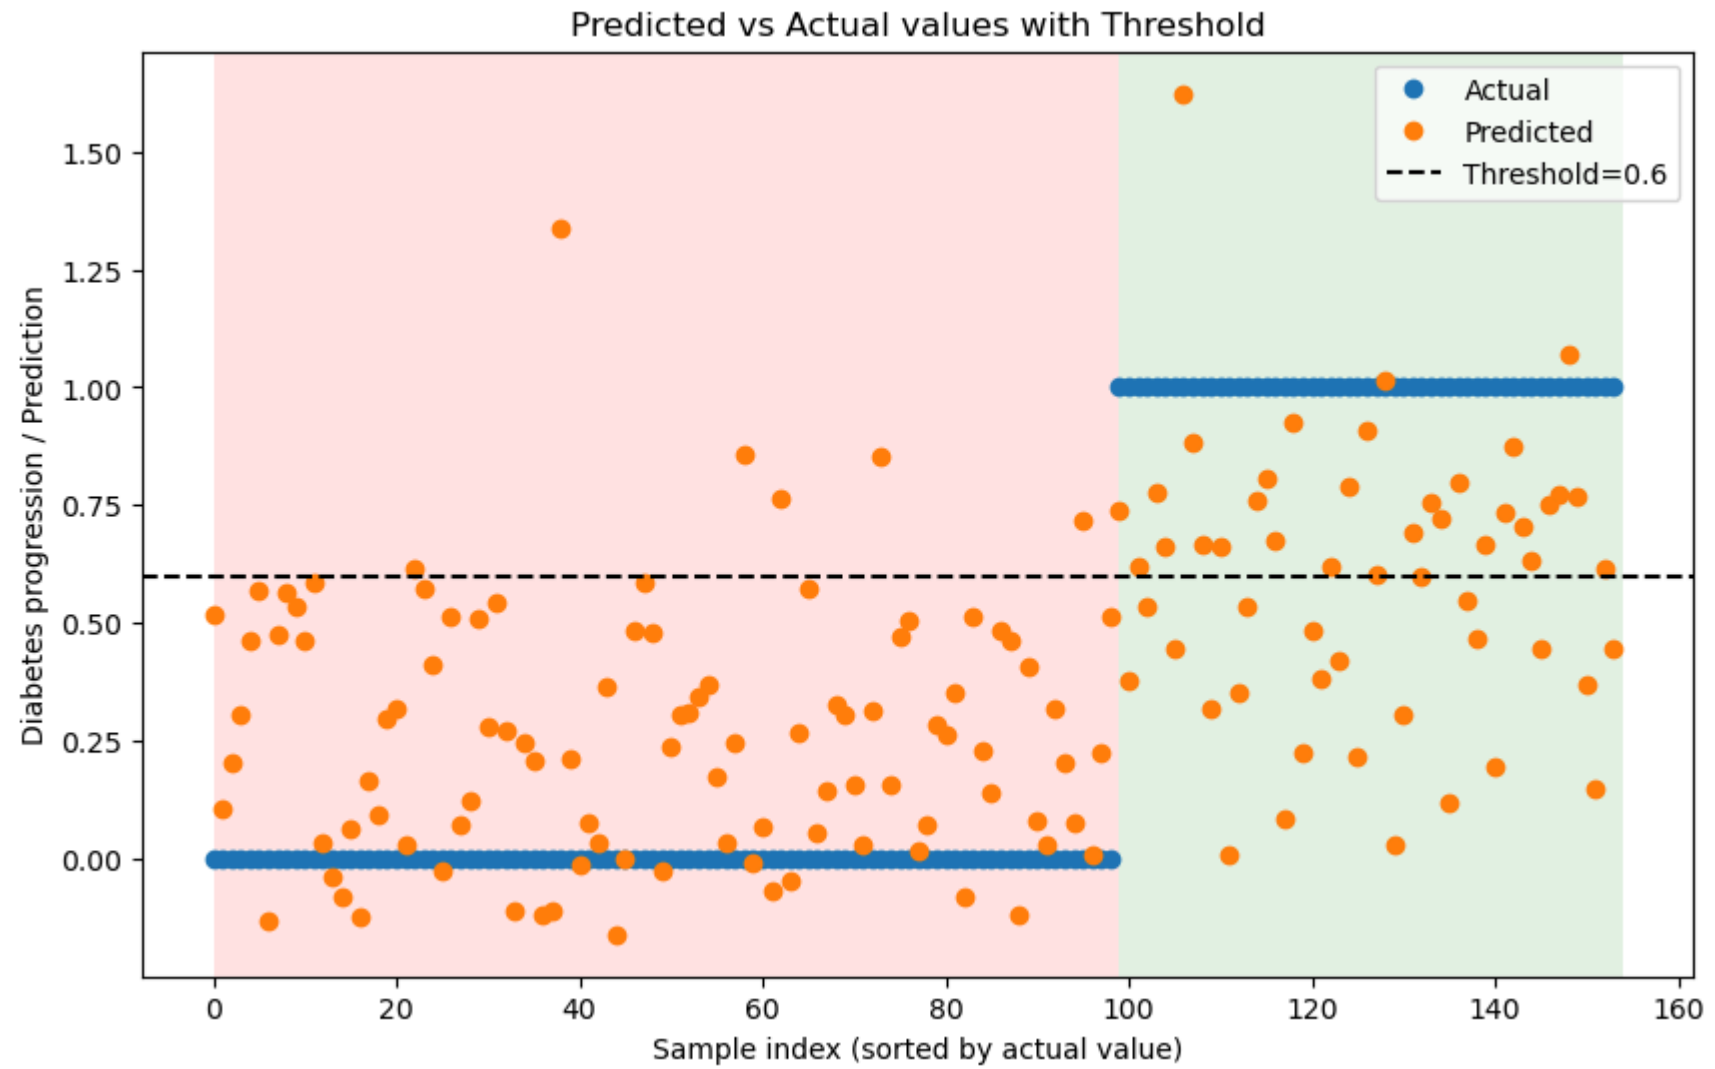
# Shade negative region (yte = 0)
plt.axvspan(0, frontier_idx, facecolor='red', alpha=0.1)

# Shade positive region (yte = 1)
plt.axvspan(frontier_idx, len(yte_sorted), facecolor='green', alpha=0.1)

# Actual and predicted points
plt.plot(range(len(yte_sorted)), yte_sorted, label="Actual", marker="o", linestyle='')
plt.plot(range(len(y_pred_sorted)), y_pred_sorted, label="Predicted", marker="o", linestyle='')

# Threshold Line
plt.axhline(y=thr, color='k', linestyle='--', label=f'Threshold={thr}')

plt.xlabel("Sample index (sorted by actual value)")
plt.ylabel("Diabetes progression / Prediction")
plt.title("Predicted vs Actual values with Threshold")
plt.legend()
plt.show()
```



```
In [9]: import numpy as np
        from sklearn.datasets import fetch_openml
        from sklearn.preprocessing import StandardScaler
        from sklearn.model_selection import train_test_split
        from sklearn.metrics import mean_squared_error, accuracy_score, classification_report, confusion_matrix
        import matplotlib.pyplot as plt
        import torch.nn.functional as F
        import torch
        import torch.nn as nn
        import torch.optim as optim
        from torch.utils.data import TensorDataset, DataLoader
        import pandas
```

Importation of Dataset 2

```
In [10]: # CHOOSE DATASET

        # Binary classification dataset
        diabetes = fetch_openml("diabetes", version = 1, as_frame=True)

        X = diabetes.data.values
        y = diabetes.target.values

        y = np.where(y == "tested_positive", 1, 0)

        print("Shape:", X.shape)

        print(diabetes.data.head(), "\n \n") # first rows of features
        print(diabetes.target.head()) # first rows of target
```

Shape: (768, 8)

	preg	plas	pres	skin	insu	mass	pedi	age
0	6	148	72	35	0	33.6	0.627	50
1	1	85	66	29	0	26.6	0.351	31
2	8	183	64	0	0	23.3	0.672	32
3	1	89	66	23	94	28.1	0.167	21
4	0	137	40	35	168	43.1	2.288	33

0 tested\_positive

1 tested\_negative

2 tested\_positive

3 tested\_negative

4 tested\_positive

Name: class, dtype: category

Categories (2, object): ['tested\_negative', 'tested\_positive']

```
In [11]: #train test splitting
test_size=0.2
Xtr, Xte, ytr, yte = train_test_split(X, y, test_size=test_size, random_state=42)
```

```
In [12]: # Standardize features
scaler=StandardScaler()
Xtr= scaler.fit_transform(Xtr)
Xte= scaler.transform(Xte)
```

A fixed seed was added to this code to ensure the reproducibility of the analysis. This allowed for manual tuning of the hyperparameters to achieve better model training.

```
In [13]: import random

seed = 42
torch.manual_seed(seed)
np.random.seed(seed)
random.seed(seed)

# Para GPU
torch.cuda.manual_seed(seed)
torch.cuda.manual_seed_all(seed)
```

```
# Tornar CUDA determinístico
torch.backends.cudnn.deterministic = True
torch.backends.cudnn.benchmark = False
```

For this model, the number of neurons per layer was increased from 64 to 100. This change was made after observing that a higher number of neurons per layer could lead to improved model performance.

```
In [14]: class MLP(nn.Module):
    def __init__(self, input_size, output_size=1, dropout_prob=0.5):
        super(MLP, self).__init__()

        a = 100

        self.fc1 = nn.Linear(input_size, a)
        self.fc2 = nn.Linear(a, a)
        self.fc3 = nn.Linear(a, a)
        self.fc4 = nn.Linear(a, a)
        self.out = nn.Linear(a, output_size)

        self.dropout = nn.Dropout(p=dropout_prob)

    def forward(self, x):
        x = F.relu(self.fc1(x))
        x = self.dropout(x)

        x = F.relu(self.fc2(x))
        x = self.dropout(x)

        x = F.relu(self.fc3(x))
        x = self.dropout(x)

        x = F.relu(self.fc4(x))
        x = self.dropout(x)

        x = self.out(x)
        return x
```



```
In [15]: num_epochs=150
         lr=0.0001
         dropout=0.1
         batch_size=128
```

This model was trained on the GPU and then transferred to the CPU for use with NumPy. Given the low number of parameters (i.e., the model's low complexity), the time required on the CPU was similar to that on the GPU.

```
In [16]: # Model, Loss, Optimizer
         device = torch.device("cuda" if torch.cuda.is_available() else "cpu")
         #device = "cpu" # force to use CPU
         print(device)
```

cuda

```
In [17]: Xtr = torch.tensor(Xtr, dtype=torch.float32).to(device)
         ytr = torch.tensor(ytr, dtype=torch.float32).to(device)
         Xte = torch.tensor(Xte, dtype=torch.float32).to(device)
         yte = torch.tensor(yte, dtype=torch.float32).to("cpu")

         # Wrap Xtr and ytr into a dataset
         train_dataset = TensorDataset(Xtr, ytr)

         # Create DataLoader
         train_dataloader = DataLoader(train_dataset, batch_size=batch_size, shuffle=True)
```

```
In [18]: model = MLP(input_size=Xtr.shape[1], dropout_prob=dropout).to(device)
         criterion = nn.BCEWithLogitsLoss() # for binary classification
         criterion = nn.MSELoss() #for regression
         optimizer = optim.Adam(model.parameters(), lr=lr)
```

The model was implemented as a fully connected neural network (MLP) with four hidden layers of 64 neurons each, using ReLU activation functions. The network takes the input features of the dataset and outputs a single value for regression (Diabetes Progression).

```
In [19]: # Training Loop
         import time
         start_time = time.time()
         for epoch in range(num_epochs):
```

```
model.train()
epoch_loss = 0.0

for batch_x, batch_y in train_dataloader:
    batch_x = batch_x.to(device)
    batch_y = batch_y.to(device)

    logits = model(batch_x)
    loss = criterion(logits, batch_y.view(-1, 1))

    optimizer.zero_grad()
    loss.backward()
    optimizer.step()

    epoch_loss += loss.item()

avg_loss = epoch_loss / len(train_dataloader)
print(f"Epoch [{epoch+1}/{num_epochs}], Loss: {avg_loss:.4f}")

end_time = time.time()
print(f"Training time: {end_time - start_time:.2f} seconds")
```

```
Epoch [1/150], Loss: 0.3047
Epoch [2/150], Loss: 0.2965
Epoch [3/150], Loss: 0.2853
Epoch [4/150], Loss: 0.2748
Epoch [5/150], Loss: 0.2666
Epoch [6/150], Loss: 0.2612
Epoch [7/150], Loss: 0.2516
Epoch [8/150], Loss: 0.2430
Epoch [9/150], Loss: 0.2342
Epoch [10/150], Loss: 0.2262
Epoch [11/150], Loss: 0.2204
Epoch [12/150], Loss: 0.2118
Epoch [13/150], Loss: 0.2021
Epoch [14/150], Loss: 0.2008
Epoch [15/150], Loss: 0.1928
Epoch [16/150], Loss: 0.1910
Epoch [17/150], Loss: 0.1854
Epoch [18/150], Loss: 0.1833
Epoch [19/150], Loss: 0.1756
Epoch [20/150], Loss: 0.1753
Epoch [21/150], Loss: 0.1742
Epoch [22/150], Loss: 0.1699
Epoch [23/150], Loss: 0.1679
Epoch [24/150], Loss: 0.1621
Epoch [25/150], Loss: 0.1665
Epoch [26/150], Loss: 0.1668
Epoch [27/150], Loss: 0.1589
Epoch [28/150], Loss: 0.1595
Epoch [29/150], Loss: 0.1589
Epoch [30/150], Loss: 0.1579
Epoch [31/150], Loss: 0.1586
Epoch [32/150], Loss: 0.1594
Epoch [33/150], Loss: 0.1581
Epoch [34/150], Loss: 0.1524
Epoch [35/150], Loss: 0.1543
Epoch [36/150], Loss: 0.1496
Epoch [37/150], Loss: 0.1515
Epoch [38/150], Loss: 0.1545
Epoch [39/150], Loss: 0.1566
Epoch [40/150], Loss: 0.1530
Epoch [41/150], Loss: 0.1583
```

Epoch [42/150], Loss: 0.1496  
Epoch [43/150], Loss: 0.1494  
Epoch [44/150], Loss: 0.1520  
Epoch [45/150], Loss: 0.1494  
Epoch [46/150], Loss: 0.1473  
Epoch [47/150], Loss: 0.1484  
Epoch [48/150], Loss: 0.1517  
Epoch [49/150], Loss: 0.1460  
Epoch [50/150], Loss: 0.1500  
Epoch [51/150], Loss: 0.1483  
Epoch [52/150], Loss: 0.1478  
Epoch [53/150], Loss: 0.1515  
Epoch [54/150], Loss: 0.1521  
Epoch [55/150], Loss: 0.1468  
Epoch [56/150], Loss: 0.1484  
Epoch [57/150], Loss: 0.1493  
Epoch [58/150], Loss: 0.1454  
Epoch [59/150], Loss: 0.1473  
Epoch [60/150], Loss: 0.1443  
Epoch [61/150], Loss: 0.1475  
Epoch [62/150], Loss: 0.1423  
Epoch [63/150], Loss: 0.1428  
Epoch [64/150], Loss: 0.1438  
Epoch [65/150], Loss: 0.1434  
Epoch [66/150], Loss: 0.1442  
Epoch [67/150], Loss: 0.1461  
Epoch [68/150], Loss: 0.1430  
Epoch [69/150], Loss: 0.1445  
Epoch [70/150], Loss: 0.1458  
Epoch [71/150], Loss: 0.1419  
Epoch [72/150], Loss: 0.1426  
Epoch [73/150], Loss: 0.1438  
Epoch [74/150], Loss: 0.1436  
Epoch [75/150], Loss: 0.1452  
Epoch [76/150], Loss: 0.1464  
Epoch [77/150], Loss: 0.1443  
Epoch [78/150], Loss: 0.1388  
Epoch [79/150], Loss: 0.1399  
Epoch [80/150], Loss: 0.1416  
Epoch [81/150], Loss: 0.1383  
Epoch [82/150], Loss: 0.1432

Epoch [83/150], Loss: 0.1387  
Epoch [84/150], Loss: 0.1417  
Epoch [85/150], Loss: 0.1406  
Epoch [86/150], Loss: 0.1373  
Epoch [87/150], Loss: 0.1430  
Epoch [88/150], Loss: 0.1435  
Epoch [89/150], Loss: 0.1341  
Epoch [90/150], Loss: 0.1422  
Epoch [91/150], Loss: 0.1343  
Epoch [92/150], Loss: 0.1393  
Epoch [93/150], Loss: 0.1401  
Epoch [94/150], Loss: 0.1424  
Epoch [95/150], Loss: 0.1357  
Epoch [96/150], Loss: 0.1378  
Epoch [97/150], Loss: 0.1424  
Epoch [98/150], Loss: 0.1363  
Epoch [99/150], Loss: 0.1360  
Epoch [100/150], Loss: 0.1367  
Epoch [101/150], Loss: 0.1409  
Epoch [102/150], Loss: 0.1403  
Epoch [103/150], Loss: 0.1374  
Epoch [104/150], Loss: 0.1308  
Epoch [105/150], Loss: 0.1308  
Epoch [106/150], Loss: 0.1387  
Epoch [107/150], Loss: 0.1351  
Epoch [108/150], Loss: 0.1383  
Epoch [109/150], Loss: 0.1396  
Epoch [110/150], Loss: 0.1402  
Epoch [111/150], Loss: 0.1359  
Epoch [112/150], Loss: 0.1393  
Epoch [113/150], Loss: 0.1298  
Epoch [114/150], Loss: 0.1328  
Epoch [115/150], Loss: 0.1302  
Epoch [116/150], Loss: 0.1371  
Epoch [117/150], Loss: 0.1325  
Epoch [118/150], Loss: 0.1379  
Epoch [119/150], Loss: 0.1357  
Epoch [120/150], Loss: 0.1342  
Epoch [121/150], Loss: 0.1400  
Epoch [122/150], Loss: 0.1363  
Epoch [123/150], Loss: 0.1359

```
Epoch [124/150], Loss: 0.1322
Epoch [125/150], Loss: 0.1306
Epoch [126/150], Loss: 0.1338
Epoch [127/150], Loss: 0.1299
Epoch [128/150], Loss: 0.1254
Epoch [129/150], Loss: 0.1297
Epoch [130/150], Loss: 0.1311
Epoch [131/150], Loss: 0.1315
Epoch [132/150], Loss: 0.1311
Epoch [133/150], Loss: 0.1289
Epoch [134/150], Loss: 0.1358
Epoch [135/150], Loss: 0.1307
Epoch [136/150], Loss: 0.1301
Epoch [137/150], Loss: 0.1312
Epoch [138/150], Loss: 0.1297
Epoch [139/150], Loss: 0.1352
Epoch [140/150], Loss: 0.1322
Epoch [141/150], Loss: 0.1337
Epoch [142/150], Loss: 0.1260
Epoch [143/150], Loss: 0.1281
Epoch [144/150], Loss: 0.1324
Epoch [145/150], Loss: 0.1312
Epoch [146/150], Loss: 0.1259
Epoch [147/150], Loss: 0.1283
Epoch [148/150], Loss: 0.1314
Epoch [149/150], Loss: 0.1234
Epoch [150/150], Loss: 0.1310
Training time: 2.52 seconds
```

```
In [20]: thr = 0.6 #threshold to tune

y_pred=model(Xte).cpu()
y_true = yte.detach().numpy()
y_pred_bin = (y_pred.detach().numpy() > thr).astype(int) # binary predictions

# Accuracy
acc = accuracy_score(y_true, y_pred_bin)
print(f'ACC:{accuracy_score(y_true,y_pred_bin)}') #classification

cm = confusion_matrix(y_true, y_pred_bin)
print("Confusion Matrix:")
```

```
print(cm)

#print(f'ACC:{mean_squared_error(yte.detach().numpy(),y_pred.detach().numpy())}') #regression
```

ACC:0.7922077922077922

Confusion Matrix:

```
[[90  9]
 [23 32]]
```

After manually tuning the hyperparameters, the best configuration obtained was: num\_epochs = 150, lr = 0.0001, dropout = 0.1, and batch\_size = 128. Dropout with a probability of 0.1 was applied after each layer to mitigate overfitting. To validate the model, both accuracy and the confusion matrix were computed. An accuracy of 79.22% was achieved, which is comparable to other models. It was observed that the model performs better when predicting negative cases compared to positive cases. This behavior is likely due to the dataset being unbalanced, highlighting the importance of data quality for achieving reliable model performance. One possible approach to address this issue is to oversample the positive class so that both classes have similar representation during training.

```
In [21]: # Convert y_pred to numpy and flatten
y_pred_np = y_pred.detach().numpy().flatten()
y_pred_bin_np = y_pred_bin
yte_np = yte.detach().numpy()

# Get indices that sort yte
sort_idx = np.argsort(yte_np)

# Sort yte and y_pred
yte_sorted = yte_np[sort_idx]
y_pred_sorted = y_pred_np[sort_idx]

# Determine the index where actual target switches from 0 to 1
frontier_idx = np.argmax(yte_sorted == 1) # first occurrence of 1

# Plot
plt.figure(figsize=(10,6))

# Shade negative region (yte = 0)
plt.axvspan(0, frontier_idx, facecolor='red', alpha=0.1)

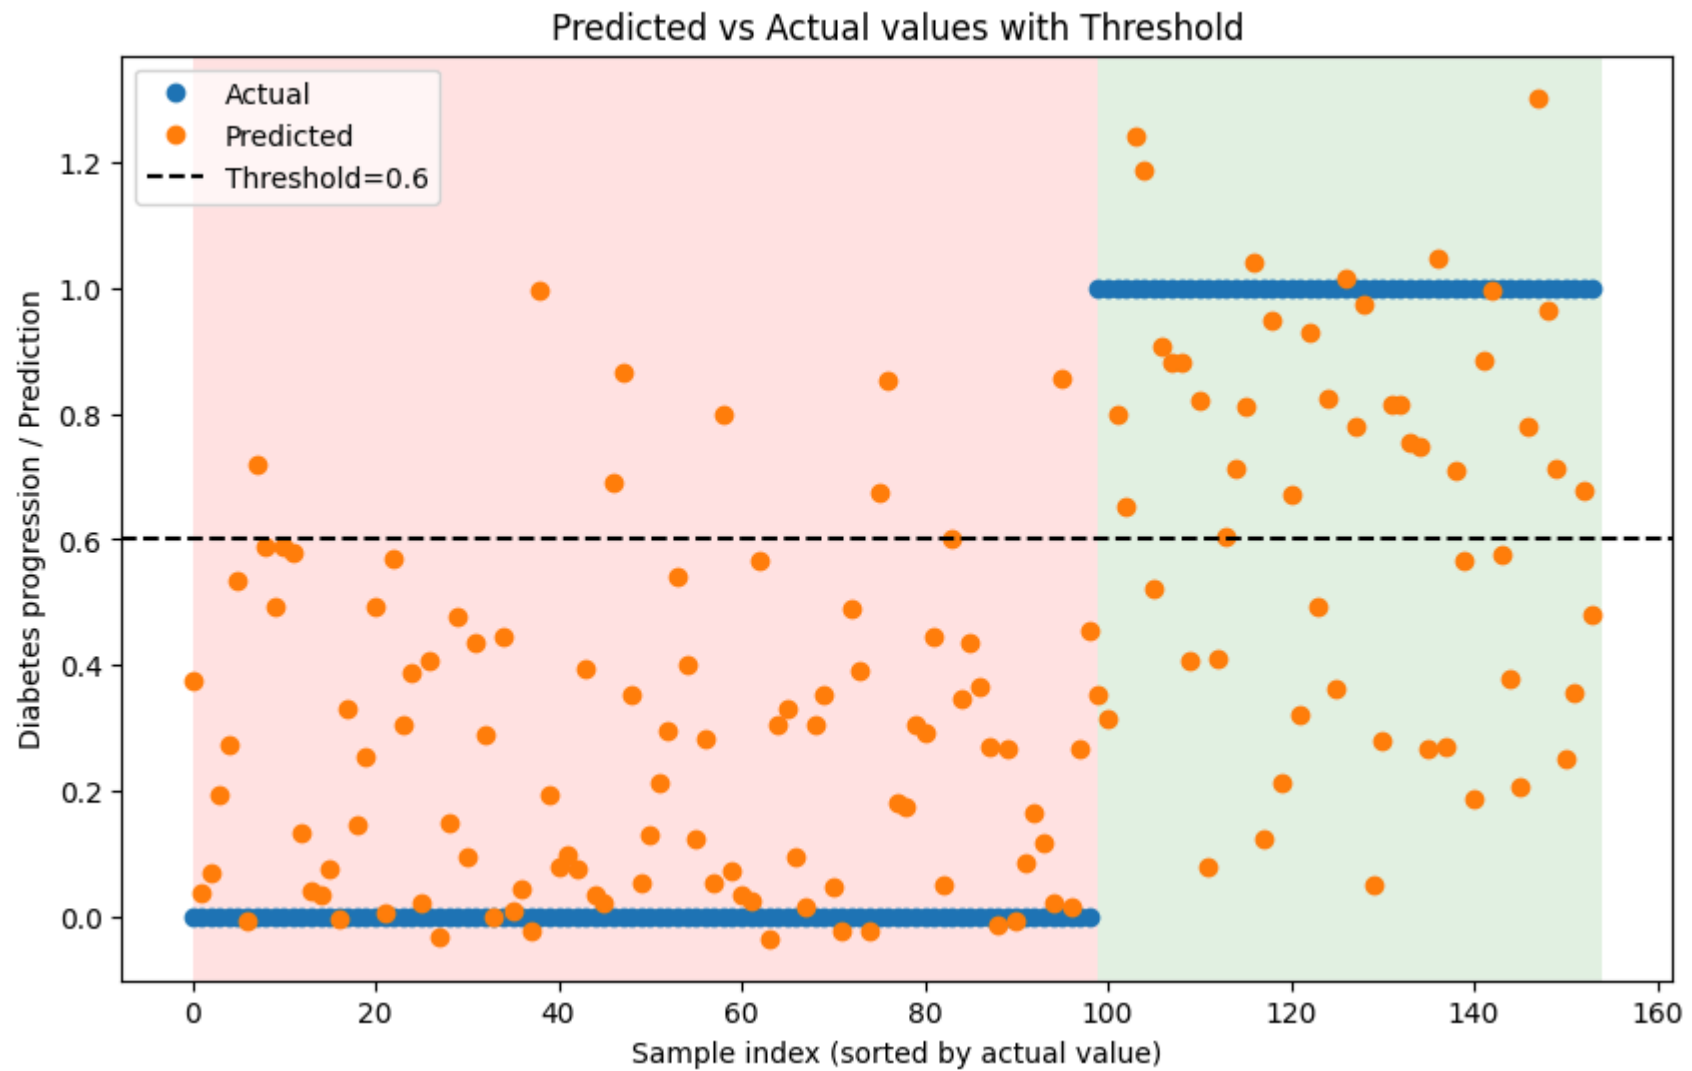
# Shade positive region (yte = 1)
plt.axvspan(frontier_idx, len(yte_sorted), facecolor='green', alpha=0.1)
```

```
# Actual and predicted points
plt.plot(range(len(yte_sorted)), yte_sorted, label="Actual", marker="o", linestyle='')
plt.plot(range(len(y_pred_sorted)), y_pred_sorted, label="Predicted", marker="o", linestyle='')

# Threshold Line
plt.axhline(y=thr, color='k', linestyle='--', label=f'Threshold={thr}')

plt.xlabel("Sample index (sorted by actual value)")
plt.ylabel("Diabetes progression / Prediction")
plt.title("Predicted vs Actual values with Threshold")
plt.legend()
plt.show()
```





In [ ]: