

```
In [1]: import numpy as np
        from sklearn import datasets
        from sklearn.preprocessing import StandardScaler
        from sklearn.model_selection import train_test_split
        from sklearn.metrics import mean_squared_error, accuracy_score, classification_report
        import matplotlib.pyplot as plt
        import torch.nn.functional as F
        import torch
        import torch.nn as nn
        import torch.optim as optim
        from torch.utils.data import TensorDataset, DataLoader
        import pandas
```

```
In [2]: # CHOOSE DATASET

        # Binary classification dataset
        diabetes = datasets.load_diabetes(as_frame=True)

        # Regression dataset
        #data = datasets.fetch_openml(name="boston", version=1, as_frame=True)

        X = diabetes.data.values
        y = diabetes.target.values

        print("Shape:", X.shape)

        print(diabetes.data.head(), "\n \n") # first rows of features
        print(diabetes.target.head()) # first rows of target
```

Shape: (442, 10)

	age	sex	bmi	bp	s1	s2	s3	\
0	0.038076	0.050680	0.061696	0.021872	-0.044223	-0.034821	-0.043401	
1	-0.001882	-0.044642	-0.051474	-0.026328	-0.008449	-0.019163	0.074412	
2	0.085299	0.050680	0.044451	-0.005670	-0.045599	-0.034194	-0.032356	
3	-0.089063	-0.044642	-0.011595	-0.036656	0.012191	0.024991	-0.036038	
4	0.005383	-0.044642	-0.036385	0.021872	0.003935	0.015596	0.008142	

	s4	s5	s6
0	-0.002592	0.019907	-0.017646
1	-0.039493	-0.068332	-0.092204
2	-0.002592	0.002861	-0.025930
3	0.034309	0.022688	-0.009362
4	-0.002592	-0.031988	-0.046641

0	151.0
1	75.0
2	141.0
3	206.0
4	135.0

Name: target, dtype: float64

```
In [3]: #train test splitting
test_size=0.2
Xtr, Xte, ytr, yte = train_test_split(X, y, test_size=test_size, random_state=42)
```

```
In [4]: # Standardize features
scaler=StandardScaler()
Xtr= scaler.fit_transform(Xtr)
Xte= scaler.transform(Xte)
```

A fixed seed was added to this code to ensure the reproducibility of the analysis. This allowed for manual tuning of the hyperparameters to achieve better model training.

```
In [5]: import random

seed = 42
torch.manual_seed(seed)
```

```
np.random.seed(seed)
random.seed(seed)

# Para GPU
torch.cuda.manual_seed(seed)
torch.cuda.manual_seed_all(seed)

# Tornar CUDA determinístico
torch.backends.cudnn.deterministic = True
torch.backends.cudnn.benchmark = False
```

```
In [6]: class MLP(nn.Module):
        def __init__(self, input_size, output_size=1, dropout_prob=0.5):
            super(MLP, self).__init__()

            self.fc1 = nn.Linear(input_size, 64)
            self.fc2 = nn.Linear(64, 64)
            self.fc3 = nn.Linear(64, 64)
            self.fc4 = nn.Linear(64, 64)
            self.out = nn.Linear(64, output_size)

            self.dropout = nn.Dropout(p=dropout_prob)

        def forward(self, x):
            x = F.relu(self.fc1(x))
            x = self.dropout(x)

            x = F.relu(self.fc2(x))
            x = self.dropout(x)

            x = F.relu(self.fc3(x))
            x = self.dropout(x)

            x = F.relu(self.fc4(x))
            x = self.dropout(x)

            x = self.out(x)
            return x
```

```
In [7]: num_epochs=190
        lr=0.01
        dropout=0.25
        batch_size=128
```

This model was trained on the GPU and then transferred to the CPU for use with NumPy. Given the low number of parameters (i.e., the model's low complexity), the time required on the CPU was similar to that on the GPU.

```
In [8]: # Model, Loss, Optimizer
        device = torch.device("cuda" if torch.cuda.is_available() else "cpu")
        #device = "cpu" # force to use CPU
        print(device)
```

cuda

```
In [9]: Xtr = torch.tensor(Xtr, dtype=torch.float32).to(device)
        ytr = torch.tensor(ytr, dtype=torch.float32).to(device)
        Xte = torch.tensor(Xte, dtype=torch.float32).to(device)
        yte = torch.tensor(yte, dtype=torch.float32).to("cpu")

        # Wrap Xtr and ytr into a dataset
        train_dataset = TensorDataset(Xtr, ytr)

        # Create DataLoader
        train_dataloader = DataLoader(train_dataset, batch_size=batch_size, shuffle=True)
```

```
In [10]: model = MLP(input_size=Xtr.shape[1], dropout_prob=dropout).to(device)
         criterion = nn.BCEWithLogitsLoss() # for binary classification
         criterion = nn.MSELoss() #for regression
         optimizer = optim.Adam(model.parameters(), lr=lr)
```

The model was implemented as a fully connected neural network (MLP) with four hidden layers of 64 neurons each, using ReLU activation functions. The network takes the input features of the dataset and outputs a single value for regression (Diabetes Progression).

```
In [11]: # Training Loop
         import time
         start_time = time.time()
         for epoch in range(num_epochs):
```

```
model.train()
epoch_loss = 0.0

for batch_x, batch_y in train_dataloader:
    batch_x = batch_x.to(device)
    batch_y = batch_y.to(device)

    logits = model(batch_x)
    loss = criterion(logits, batch_y.view(-1, 1))

    optimizer.zero_grad()
    loss.backward()
    optimizer.step()

    epoch_loss += loss.item()

avg_loss = epoch_loss / len(train_dataloader)
print(f"Epoch [{epoch+1}/{num_epochs}], Loss: {avg_loss:.4f}")

end_time = time.time()
print(f"Training time: {end_time - start_time:.2f} seconds")
```

Epoch [1/190], Loss: 29648.7161  
Epoch [2/190], Loss: 27738.2480  
Epoch [3/190], Loss: 19526.0052  
Epoch [4/190], Loss: 10916.4779  
Epoch [5/190], Loss: 7908.8825  
Epoch [6/190], Loss: 6315.4227  
Epoch [7/190], Loss: 6881.3011  
Epoch [8/190], Loss: 5394.8262  
Epoch [9/190], Loss: 5833.4985  
Epoch [10/190], Loss: 4486.8579  
Epoch [11/190], Loss: 4801.0705  
Epoch [12/190], Loss: 4380.2632  
Epoch [13/190], Loss: 4705.3464  
Epoch [14/190], Loss: 4574.8371  
Epoch [15/190], Loss: 4115.9676  
Epoch [16/190], Loss: 4263.3736  
Epoch [17/190], Loss: 4220.0199  
Epoch [18/190], Loss: 3679.3091  
Epoch [19/190], Loss: 3745.2568  
Epoch [20/190], Loss: 3896.5340  
Epoch [21/190], Loss: 4014.5900  
Epoch [22/190], Loss: 3893.0127  
Epoch [23/190], Loss: 3573.5227  
Epoch [24/190], Loss: 3891.4432  
Epoch [25/190], Loss: 3775.5807  
Epoch [26/190], Loss: 3753.8115  
Epoch [27/190], Loss: 3697.5588  
Epoch [28/190], Loss: 3638.7832  
Epoch [29/190], Loss: 4002.4415  
Epoch [30/190], Loss: 3921.0071  
Epoch [31/190], Loss: 3542.5956  
Epoch [32/190], Loss: 3535.1024  
Epoch [33/190], Loss: 3598.4771  
Epoch [34/190], Loss: 3983.6483  
Epoch [35/190], Loss: 3522.1357  
Epoch [36/190], Loss: 3819.0868  
Epoch [37/190], Loss: 3664.6637  
Epoch [38/190], Loss: 3815.6850  
Epoch [39/190], Loss: 3901.3689  
Epoch [40/190], Loss: 3601.0779  
Epoch [41/190], Loss: 3759.6592

Epoch [42/190], Loss: 3391.2233  
Epoch [43/190], Loss: 3242.0986  
Epoch [44/190], Loss: 3577.6099  
Epoch [45/190], Loss: 3480.0977  
Epoch [46/190], Loss: 3555.8261  
Epoch [47/190], Loss: 3476.3621  
Epoch [48/190], Loss: 3268.5854  
Epoch [49/190], Loss: 3618.4499  
Epoch [50/190], Loss: 3584.3310  
Epoch [51/190], Loss: 3675.5367  
Epoch [52/190], Loss: 3722.2685  
Epoch [53/190], Loss: 3227.3150  
Epoch [54/190], Loss: 2974.1118  
Epoch [55/190], Loss: 3562.6582  
Epoch [56/190], Loss: 3198.2245  
Epoch [57/190], Loss: 3238.8750  
Epoch [58/190], Loss: 3376.6782  
Epoch [59/190], Loss: 3346.6515  
Epoch [60/190], Loss: 3468.1538  
Epoch [61/190], Loss: 3152.5792  
Epoch [62/190], Loss: 3484.5428  
Epoch [63/190], Loss: 3397.3236  
Epoch [64/190], Loss: 3454.3319  
Epoch [65/190], Loss: 3955.6242  
Epoch [66/190], Loss: 3623.3734  
Epoch [67/190], Loss: 3310.0218  
Epoch [68/190], Loss: 3456.5499  
Epoch [69/190], Loss: 3696.4705  
Epoch [70/190], Loss: 3757.5662  
Epoch [71/190], Loss: 3274.0444  
Epoch [72/190], Loss: 3489.4048  
Epoch [73/190], Loss: 3018.0824  
Epoch [74/190], Loss: 3315.5653  
Epoch [75/190], Loss: 3203.5889  
Epoch [76/190], Loss: 3338.0278  
Epoch [77/190], Loss: 3350.0305  
Epoch [78/190], Loss: 3378.0791  
Epoch [79/190], Loss: 3423.2143  
Epoch [80/190], Loss: 3527.6632  
Epoch [81/190], Loss: 3268.2509  
Epoch [82/190], Loss: 3227.4768

Epoch [83/190], Loss: 3589.3934  
Epoch [84/190], Loss: 3299.7594  
Epoch [85/190], Loss: 3385.5459  
Epoch [86/190], Loss: 3498.3327  
Epoch [87/190], Loss: 3477.1771  
Epoch [88/190], Loss: 3174.3293  
Epoch [89/190], Loss: 3714.6216  
Epoch [90/190], Loss: 3343.3033  
Epoch [91/190], Loss: 3414.0710  
Epoch [92/190], Loss: 3248.5864  
Epoch [93/190], Loss: 3488.2712  
Epoch [94/190], Loss: 3373.2797  
Epoch [95/190], Loss: 3372.9124  
Epoch [96/190], Loss: 3554.2501  
Epoch [97/190], Loss: 3368.5915  
Epoch [98/190], Loss: 3215.7687  
Epoch [99/190], Loss: 3260.8147  
Epoch [100/190], Loss: 3192.8174  
Epoch [101/190], Loss: 3423.2056  
Epoch [102/190], Loss: 2991.1534  
Epoch [103/190], Loss: 3152.7769  
Epoch [104/190], Loss: 3303.9213  
Epoch [105/190], Loss: 3282.2486  
Epoch [106/190], Loss: 3331.6805  
Epoch [107/190], Loss: 3279.8656  
Epoch [108/190], Loss: 3262.2849  
Epoch [109/190], Loss: 3125.2458  
Epoch [110/190], Loss: 3407.5113  
Epoch [111/190], Loss: 3498.1632  
Epoch [112/190], Loss: 3553.5989  
Epoch [113/190], Loss: 3089.8211  
Epoch [114/190], Loss: 3075.9001  
Epoch [115/190], Loss: 3539.4534  
Epoch [116/190], Loss: 2961.9173  
Epoch [117/190], Loss: 3655.1819  
Epoch [118/190], Loss: 3049.0596  
Epoch [119/190], Loss: 3615.5492  
Epoch [120/190], Loss: 3256.5053  
Epoch [121/190], Loss: 3308.7235  
Epoch [122/190], Loss: 3240.1512  
Epoch [123/190], Loss: 3193.3521



Epoch [124/190], Loss: 3216.8752  
Epoch [125/190], Loss: 3073.2205  
Epoch [126/190], Loss: 3247.6636  
Epoch [127/190], Loss: 3242.5971  
Epoch [128/190], Loss: 3164.6580  
Epoch [129/190], Loss: 3109.4061  
Epoch [130/190], Loss: 3200.9823  
Epoch [131/190], Loss: 2903.9763  
Epoch [132/190], Loss: 3002.0817  
Epoch [133/190], Loss: 3106.7635  
Epoch [134/190], Loss: 3022.0481  
Epoch [135/190], Loss: 2955.9606  
Epoch [136/190], Loss: 2726.1379  
Epoch [137/190], Loss: 2718.2030  
Epoch [138/190], Loss: 2962.2109  
Epoch [139/190], Loss: 2859.1110  
Epoch [140/190], Loss: 2965.7663  
Epoch [141/190], Loss: 2804.8240  
Epoch [142/190], Loss: 3106.8247  
Epoch [143/190], Loss: 2979.2811  
Epoch [144/190], Loss: 3120.4154  
Epoch [145/190], Loss: 3010.3355  
Epoch [146/190], Loss: 2924.7756  
Epoch [147/190], Loss: 3123.8313  
Epoch [148/190], Loss: 3109.5750  
Epoch [149/190], Loss: 3098.7642  
Epoch [150/190], Loss: 2857.3700  
Epoch [151/190], Loss: 2996.3384  
Epoch [152/190], Loss: 3129.6170  
Epoch [153/190], Loss: 2759.2476  
Epoch [154/190], Loss: 3000.5859  
Epoch [155/190], Loss: 2925.3764  
Epoch [156/190], Loss: 2861.1976  
Epoch [157/190], Loss: 2986.6821  
Epoch [158/190], Loss: 2743.7799  
Epoch [159/190], Loss: 3126.5764  
Epoch [160/190], Loss: 2914.5562  
Epoch [161/190], Loss: 2880.0532  
Epoch [162/190], Loss: 2857.0591  
Epoch [163/190], Loss: 2817.7529  
Epoch [164/190], Loss: 2701.1991

```
Epoch [165/190], Loss: 3057.9538
Epoch [166/190], Loss: 2827.3732
Epoch [167/190], Loss: 2802.8832
Epoch [168/190], Loss: 3131.7845
Epoch [169/190], Loss: 3075.6080
Epoch [170/190], Loss: 2931.3674
Epoch [171/190], Loss: 2838.9221
Epoch [172/190], Loss: 3041.7693
Epoch [173/190], Loss: 3094.6465
Epoch [174/190], Loss: 3209.8215
Epoch [175/190], Loss: 3203.1895
Epoch [176/190], Loss: 3175.8590
Epoch [177/190], Loss: 2773.8187
Epoch [178/190], Loss: 3101.5061
Epoch [179/190], Loss: 3159.9993
Epoch [180/190], Loss: 2911.1960
Epoch [181/190], Loss: 2810.7509
Epoch [182/190], Loss: 2745.5351
Epoch [183/190], Loss: 2873.0893
Epoch [184/190], Loss: 3086.1829
Epoch [185/190], Loss: 3048.7624
Epoch [186/190], Loss: 2893.8034
Epoch [187/190], Loss: 2915.0261
Epoch [188/190], Loss: 2572.4602
Epoch [189/190], Loss: 2886.5369
Epoch [190/190], Loss: 2821.7037
Training time: 1.99 seconds
```

```
In [12]: y_pred=model(Xte).cpu() # só nesta altura volta ao CPU
          #print(f'ACC:{accuracy_score(yte.detach().numpy(),y_pred.detach().numpy())>0.5}') #classification

          print(f'MSE:{mean_squared_error(yte.detach().numpy(),y_pred.detach().numpy())}') #regression
```

MSE: 2814.833251953125

After tuning the hyperparameters, the best configuration was: num\_epochs = 200, lr = 0.02, dropout = 0.2, and batch\_size = 128. A dropout with a probability of 0.2 was applied after each layer to mitigate overfitting. In order to validate the model the MSE was computed. The model was evaluated using the Mean Squared Error (MSE), which resulted in a high value of 2815. This indicates that, despite the training and parameter optimization, the current approach is insufficient to accurately predict Diabetes Progression. The high error may be attributed to

the limited size of the dataset, the low number of features. In comparison with previous approaches, such as ANFIS or TSK models, this approach exhibited the worst performance.

```
In [13]: # Plot predictions vs actual
# Converter y_pred para numpy e flatten
y_pred_np = y_pred.detach().numpy().flatten()
yte_np = yte.detach().numpy()

# Obter índices que ordenam yte
sort_idx = np.argsort(yte_np)

# Ordenar yte e y_pred segundo esses índices
yte_sorted = yte_np[sort_idx]
y_pred_sorted = y_pred_np[sort_idx]

# Plot
plt.figure(figsize=(10,6))
plt.plot(range(len(yte_sorted)), yte_sorted, label="Actual", marker="o", linestyle='')
plt.plot(range(len(y_pred_sorted)), y_pred_sorted, label="Predicted", marker="x", linestyle='')

plt.xlabel("Sample index (sorted by actual value)")
plt.ylabel("Diabetes progression")
plt.title("Predicted vs Actual values on Diabetes dataset (sorted)")
plt.legend()
plt.show()
```

